

Import required libraries

```
In [1]: import os
import numpy as np
import pandas as pd

import sklearn.linear_model
import sklearn.tree
import sklearn.metrics

from matplotlib import pyplot as plt
import seaborn as sns
```

1.1: Compute true/false positives/negatives.

```
In [2]: def calc_TP_TN_FP_FN(ytrue_N, yhat_N):
        ''' Compute counts of four possible outcomes of a binary classifier f

        Args
        ----
        ytrue_N : 1D array of floats
            Each entry represents the binary value (0 or 1) of 'true' label o
            One entry per example in current dataset
        yhat_N : 1D array of floats
            Each entry represents a predicted binary value (either 0 or 1).
            One entry per example in current dataset.
            Needs to be same size as ytrue_N.

        Returns
        -----
        TP : float
            Number of true positives
        TN : float
            Number of true negatives
        FP : float
            Number of false positives
        FN : float
            Number of false negatives
        ...

        TP = 0.0
        TN = 0.0
        FP = 0.0
        FN = 0.0

        for i in range(len(yhat_N)):
            if ytrue_N[i]==yhat_N[i]==1:
                TP +=1
            if ytrue_N[i]==yhat_N[i]==0:
                TN += 1
            if yhat_N[i]==1 and ytrue_N[i]==0:
                FP +=1
            if yhat_N[i]==0 and ytrue_N[i]==1:
                FN +=1

        return TP, TN, FP, FN
```

Supplied functions for later use

```
In [3]: def calc_perf_metrics_for_threshold(ytrue_N, yprobal_N, thresh):
        ''' Compute performance metrics for a given probabilistic classifier
        ...

        tp, tn, fp, fn = calc_TP_TN_FP_FN(ytrue_N, yprobal_N >= thresh)
        ## Compute ACC, TPR, TNR, etc.
```

```

    acc = (tp + tn) / float(tp + tn + fp + fn + 1e-10)
    tpr = tp / float(tp + fn + 1e-10)
    tnr = tn / float(fp + tn + 1e-10)
    ppv = tp / float(tp + fp + 1e-10)
    npv = tn / float(tn + fn + 1e-10)

    return acc, tpr, tnr, ppv, npv

def print_perf_metrics_for_threshold(ytrue_N, yprobal_N, thresh):
    ''' Pretty print perf. metrics for a given probabilistic classifier a
    '''
    acc, tpr, tnr, ppv, npv = calc_perf_metrics_for_threshold(ytrue_N, yp

    ## Pretty print the results
    print("%.3f ACC" % acc)
    print("%.3f TPR" % tpr)
    print("%.3f TNR" % tnr)
    print("%.3f PPV" % ppv)
    print("%.3f NPV" % npv)

def calc_confusion_matrix_for_threshold(ytrue_N, yprobal_N, thresh):
    ''' Compute the confusion matrix for a given probabilistic classifier

    Args
    ----
    ytrue_N : 1D array of floats
        Each entry represents the binary value (0 or 1) of 'true' label o
        One entry per example in current dataset
    yprobal_N : 1D array of floats
        Each entry represents a probability (between 0 and 1) that correc
        One entry per example in current dataset
        Needs to be same size as ytrue_N
    thresh : float
        Scalar threshold for converting probabilities into hard decisions
        Calls an example "positive" if yprobal >= thresh

    Returns
    ----
    cm_df : Pandas DataFrame
        Can be printed like print(cm_df) to easily display results
    '''
    cm = sklearn.metrics.confusion_matrix(ytrue_N, yprobal_N >= thresh)
    cm_df = pd.DataFrame(data=cm, columns=[0, 1], index=[0, 1])
    cm_df.columns.name = 'Predicted'
    cm_df.index.name = 'True'
    return cm_df

```

```

In [4]: def compute_perf_metrics_across_thresholds(ytrue_N, yprobal_N, thresh_gri
    ''' Compute common binary classifier performance metrics across many

```

compute common binary classifier performance metrics across many

If no array of thresholds is provided, will use all 'unique' values in the yprobal_N array to define all possible thresholds with difference

Args

ytrue_N : 1D array of floats

Each entry represents the binary value (0 or 1) of 'true' label of
One entry per example in current dataset

yprobal_N : 1D array of floats

Each entry represents a probability (between 0 and 1) that correct
One entry per example in current dataset

Returns

thresh_grid : 1D array of floats

One entry for each possible threshold

perf_dict : dict, with key, value pairs:

- * 'acc' : 1D array of accuracy values (one per threshold)
- * 'ppv' : 1D array of positive predictive values (one per threshold)
- * 'npv' : 1D array of negative predictive values (one per threshold)
- * 'tpr' : 1D array of true positive rates (one per threshold)
- * 'tnr' : 1D array of true negative rates (one per threshold)

...

if thresh_grid **is** **None**:

bin_edges = np.linspace(0, 1.001, 21)

thresh_grid = np.sort(np.hstack([bin_edges, np.unique(yprobal_N)]))

tpr_grid = np.zeros_like(thresh_grid)

tnr_grid = np.zeros_like(thresh_grid)

ppv_grid = np.zeros_like(thresh_grid)

npv_grid = np.zeros_like(thresh_grid)

acc_grid = np.zeros_like(thresh_grid)

for tt, thresh **in** enumerate(thresh_grid):

Apply specific threshold to convert probas into hard binary values

Then count number of true positives, true negatives, etc.

Then compute metrics like accuracy and true positive rate

acc, tpr, tnr, ppv, npv = calc_perf_metrics_for_threshold(ytrue_N,

acc_grid[tt] = acc

tpr_grid[tt] = tpr

tnr_grid[tt] = tnr

ppv_grid[tt] = ppv

npv_grid[tt] = npv

return thresh_grid, dict(

acc=acc_grid,

tpr=tpr_grid,

tnr=tnr_grid,

ppv=ppv_grid,

npv=npv_grid)

def make_plot_perf_vs_threshold(ytrue_N, yprobal_N, bin_edges=np.linspace

```
''' Make pretty plot of binary classifier performance as threshold in

Produces a plot with 3 rows:
* top row: hist of predicted probabilities for negative examples (sha
* middle row: hist of predicted probabilities for positive examples (
* bottom row: line plots of metrics that require hard decisions (ACC,
'''

fig, axes = plt.subplots(nrows=3, ncols=1, figsize=(12, 8))
sns.distplot(
    yprobal_N[ytrue_N == 0],
    color='r', bins=bin_edges, kde=False, rug=True, ax=axes[0]);
sns.distplot(
    yprobal_N[ytrue_N == 1],
    color='b', bins=bin_edges, kde=False, rug=True, ax=axes[1]);

thresh_grid, perf_grid = compute_perf_metrics_across_thresholds(ytrue
axes[2].plot(thresh_grid, perf_grid['acc'], 'k-', label='accuracy')
axes[2].plot(thresh_grid, perf_grid['tpr'], 'b-', label='TPR (recall/
axes[2].plot(thresh_grid, perf_grid['tnr'], 'g-', label='TNR (specifi
axes[2].plot(thresh_grid, perf_grid['ppv'], 'c-', label='PPV (precisi
axes[2].plot(thresh_grid, perf_grid['npv'], 'm-', label='NPV')

axes[2].legend()
axes[2].set_ylim([0, 1])
```

1.2: Load the dataset.

```
In [5]: # Load 3 feature version of x arrays
x_tr_M3 = np.loadtxt('./data_cancer/x_train.csv', delimiter=',', skiprows
x_va_N3 = np.loadtxt('./data_cancer/x_valid.csv', delimiter=',', skiprows
x_te_N3 = np.loadtxt('./data_cancer/x_test.csv', delimiter=',', skiprows=

# 2 feature version of x arrays
x_tr_M2 = x_tr_M3[:, :2].copy()
x_va_N2 = x_va_N3[:, :2].copy()
x_te_N2 = x_te_N3[:, :2].copy()
```

```
In [6]: y_tr_M = np.loadtxt('./data_cancer/y_train.csv', delimiter=',', skiprows=
y_va_N = np.loadtxt('./data_cancer/y_valid.csv', delimiter=',', skiprows=
y_te_N = np.loadtxt('./data_cancer/y_test.csv', delimiter=',', skiprows=1
```

1.3: The predict-0-always baseline

(a) Compute the accuracy of the always-0 classifier.

compute and print the accuracy of the always-0 classifier on validation and test outputs.

```
In [7]: yhat_N=np.zeros(180,)
        TP1,TN1,FP, FN=calc_TP_TN_FP_FN(y_va_N, yhat_N)
        acc1 = (TP1+TN1)/len(y_va_N)
        TP2,TN2,FP, FN=calc_TP_TN_FP_FN(y_te_N, yhat_N)
        acc2 = (TP2+TN2)/len(y_te_N)

        print("Always-0: accuracy on VALID: %.3f" % acc1)
        print("Always-0: accuracy on TEST : %.3f" % acc2)
```

Always-0: accuracy on VALID: 0.861

Always-0: accuracy on TEST : 0.861

(b) Print a confusion matrix for the always-0 classifier.

generate a confusion matrix for the always-0 classifier on the validation set.

```
In [8]: yprobal_N=np.zeros(180,)
        print(calc_confusion_matrix_for_threshold(y_va_N, yprobal_N, thresh=1.0))
```

Predicted	0	1
True		
0	155	0
1	25	0

The accuracy of the always-0 classifier seems decent. However, I wouldn't use this classifier in this problem, because we should care about picking out those people who actually have cancer.

1.4: Logistic Regression

(a) Create a set of LogisticRegression models.

Each model will use a different control parameter, C , and each will be fit to 2-feature data. Probabilistic predictions will be made on both training set and validation set inputs, and logistic-loss for each will be recorded.

```

In [9]: tr_loss_list = list()
        va_loss_list = list()

        C_grid = np.logspace(-9, 6, 31)

        for C in C_grid:

            logreg = sklearn.linear_model.LogisticRegression(penalty='l2', C=C, s
            logreg.fit(x_tr_M2, y_tr_M) # fit model
            # make class predictions for the testing set
            y_pred_proba = logreg.predict_proba(x_tr_M2)[: ,1] # convention
            loss = sklearn.metrics.log_loss(y_tr_M, y_pred_proba)
            tr_loss_list.append(loss)

            y_pred_proba2 = logreg.predict_proba(x_va_N2)[: ,1] # convention
            loss2 = sklearn.metrics.log_loss(y_va_N, y_pred_proba2)
            va_loss_list.append(loss2)

        best_va = np.argmin(va_loss_list)
        best_va_c = C_grid[best_va]
        print(best_va_c)
        bestmodel_va = sklearn.linear_model.LogisticRegression(penalty='l2', C=be
        bestmodel_va

```

31.622776601683793


```

Out[9]: LogisticRegression(C=31.622776601683793, class_weight=None, dual=False
,
                        fit_intercept=True, intercept_scaling=1, l1_ratio=N
one,
                        max_iter=100, multi_class='warn', n_jobs=None, pena
lty='l2',
                        random_state=0, solver='liblinear', tol=0.0001, ver
bose=0,
                        warm_start=False)

```

```
In [10]: C_grid  
plt.plot(C_grid)
```

```
Out[10]: [<matplotlib.lines.Line2D at 0x20be5a93908>]
```



Plot logistic loss (y-axis) vs. C (x-axis) on the training set and validation set.


```

In [11]: plt.xlabel('log10(C)');
plt.ylabel('logistic loss');
plt.ylim([0.0, 0.7]);

plt.plot(np.log10(C_grid),tr_loss_list,marker='o')
plt.plot(np.log10(C_grid),va_loss_list,c="red",marker='+')


min_logloss = sorted(va_loss_list)[0]

print("Best C-value for LR with 2-feature data: %.3f" % best_va_c)
print("Validation set log-loss at best C-value: %.4f" % min_logloss)

```

Best C-value for LR with 2-feature data: 31.623


Validation set log-loss at best C-value: 0.3549



(b) Plot the performance of the predictions made by the best classifier from step (a) on the validation set.

```
In [12]: bestmodel_va.fit(x_tr_M2, y_tr_M)
y_pred = bestmodel_va.predict_proba(x_va_N2)[: ,1]

make_plot_perf_vs_threshold(y_va_N, y_pred, bin_edges=np.linspace(0, 1, 2
```



(c) Model fitting with 3-feature data

```

In [13]: tr_loss_list3 = list()
         va_loss_list3 = list()

         C_grid = np.logspace(-9, 6, 31)

         # Build and evaluate model for each value C
         for C in C_grid:

             logreg = sklearn.linear_model.LogisticRegression(penalty='l2', C=C, s
             logreg.fit(x_tr_M3, y_tr_M) # fit model
             y_pred_proba = logreg.predict_proba(x_tr_M3)[: ,1] # convention
             loss = sklearn.metrics.log_loss(y_tr_M, y_pred_proba)
             tr_loss_list3.append(loss)

             # make class predictions for the testing set
             y_pred_proba2 = logreg.predict_proba(x_va_N3)[: ,1] # convention
             loss2 = sklearn.metrics.log_loss(y_va_N, y_pred_proba2)
             va_loss_list3.append(loss2)

         # bestmodel_tr = sklearn.linear_model.LogisticRegression(penalty='l2', C=

         best_va3 = np.argmin(va_loss_list3)
         best_va_c3 = C_grid[best_va3]
         bestmodel_va3 = sklearn.linear_model.LogisticRegression(penalty='l2', C=b
         bestmodel_va3

```

```

Out[13]: LogisticRegression(C=1000000.0, class_weight=None, dual=False,
                             fit_intercept=True, intercept_scaling=1, l1_ratio=N
one,
                             max_iter=100, multi_class='warn', n_jobs=None, pena
lty='l2',
                             random_state=0, solver='liblinear', tol=0.0001, ver
bose=0,
                             warm_start=False)

```

Plot logistic loss (y-axis) vs. C (x-axis) for the 3-feature classifiers on the training set and validation set.

```
In [14]: plt.xlabel('log10(C)');
plt.ylabel('logistic loss');
plt.ylim([0.0, 0.7]);


plt.plot(np.log10(C_grid),tr_loss_list3,marker='o')
plt.plot(np.log10(C_grid),va_loss_list3,c="red",marker='+')

min_logloss = sorted(va_loss_list)[0]

print("Best C-value for LR with 3-feature data: %.3f" % best_va3)
print("Validation set log-loss at best C-value: %.4f" % min_logloss)
```

Best C-value for LR with 3-feature data: 30.000


Validation set log-loss at best C-value: 0.3549



Plot the performance of the predictions made by the best 3-valued classifier on the validation set.

In [15]:

```
bestmodel_va3.fit(x_tr_M3, y_tr_M)
y_pred = bestmodel_va3.predict_proba(x_va_N3)[: ,1]
make_plot_perf_vs_threshold(y_va_N, y_pred, bin_edges=np.linspace(0, 1, 2
```




1.5: ROC Curves

These curves allow us to compare model performance in terms of trade-offs between false positive and true positive results.

(a) Plot ROC curves on the validation set.

In [16]:


```
fpr1, tpr1, thresholds1 =sklearn.metrics.roc_curve(y_va_N, bestmodel_va.p  
fpr2, tpr2, thresholds2 =sklearn.metrics.roc_curve(y_va_N, bestmodel_va3.  
plt.plot(fpr1, tpr1)  
plt.plot(fpr2, tpr2,c="red")  
  
plt.ylim([0, 1]);  
plt.xlabel("False Postive rate (FPR = 1 - TPR)");  
plt.ylabel("True Postive rate (TPR)");
```



(b) Plot ROC curves on the test set.

```
In [17]: fpr1, tpr1, thresholds1 =sklearn.metrics.roc_curve(y_te_N, bestmodel_va.p
fpr2, tpr2, thresholds2 =sklearn.metrics.roc_curve(y_te_N, bestmodel_va3.
plt.plot(fpr1, tpr1)
plt.plot(fpr2, tpr2,c="red")

plt.ylim([0, 1]);
plt.xlabel("False Postive rate (FPR = 1 - TPR)");
plt.ylabel("True Postive rate (TPR)");
```



A perfect classifier would give us 100% success for true positives, with a 0% rate of false ones, so the ROC curve that skewered left-top indicates that this model performance better. From our plots, we can know that the model with 3-feature dominates 2-feature one. Obviously, the area under the ROC will also be greater. Thus, it tells that the last feature provides more information and we should use 3 features to build the model.

1.6: Selecting a decision threshold

(a) Using default 0.5 threshold.

Generate a confusion matrix for the best 3-feature logistic model on the test set, using threshold 0.5 to start.

```
In [18]: best_thr = 0.5

print("ON THE VALIDATION SET:")
print("Chosen best thr = %.4f" % best_thr)
print("")
print("ON THE TEST SET:")

print(calc_confusion_matrix_for_threshold(y_te_N, bestmodel_va3.predict_p
print(""))

print_perf_metrics_for_threshold(y_te_N, bestmodel_va3.predict_proba(x_te
```

```
ON THE VALIDATION SET:
Chosen best thr = 0.5000
```

```
ON THE TEST SET:
Predicted    0    1
True
0           152    3
1           15   10
```

```
0.900 ACC
0.400 TPR
0.981 TNR
0.769 PPV
0.910 NPV
```

(b) Pick a threshold to maximize TPR, while ensuring PPV \geq 0.98.

After finding the best threshold on the validation set, plot its confusion matrix and print its various performance metrics, for the test set.


```
In [19]: thresh_grid, perf_grid = compute_perf_metrics_across_thresholds(y_va_N, b

index = []
for i in range(len(thresh_grid)):
    if perf_grid["ppv"][i] >= 0.98:
        index.append(i)
max_index = index[np.argmax(perf_grid["tpr"][index])]
best_thr = thresh_grid[max_index]

print("ON THE VALIDATION SET:")
print("Chosen best thr = %.4f" % best_thr)
print("")
print("ON THE TEST SET:")

print(calc_confusion_matrix_for_threshold(y_te_N, bestmodel_va3.predict_p
print(""))

print_perf_metrics_for_threshold(y_te_N, bestmodel_va3.predict_proba(x_te
```

ON THE VALIDATION SET:
Chosen best thr = 0.6290

ON THE TEST SET:

Predicted	0	1
True		
0	155	0
1	20	5

0.889 ACC
0.200 TPR
1.000 TNR
1.000 PPV
0.886 NPV

(c) Pick a threshold to maximize PPV, while ensuring TPR \geq 0.98.

After finding the best threshold on the validation set, plot its confusion matrix and print its various performance metrics, for the test set.

```
In [20]: thresh_grid, perf_grid = compute_perf_metrics_across_thresholds(y_va_N, b

index = []
for i in range(len(thresh_grid)):
    if perf_grid["tpr"][i] >= 0.98:
        index.append(i)
max_index = index[np.argmax(perf_grid["ppv"][index])]
best_thr = thresh_grid[max_index]

print("ON THE VALIDATION SET:")
print("Chosen best thr = %.4f" % best_thr)
print("")
print("ON THE TEST SET:")

print(calc_confusion_matrix_for_threshold(y_te_N, bestmodel_va3.predict_p
print(""))

print_perf_metrics_for_threshold(y_te_N, bestmodel_va3.predict_proba(x_te
```

ON THE VALIDATION SET:
Chosen best thr = 0.0300

ON THE TEST SET:

Predicted	0	1
True		
0	57	98
1	0	25

0.456 ACC
1.000 TPR
0.368 TNR
0.203 PPV
1.000 NPV

The threshold to maximize TPR, while ensuring PPV ≥ 0.98 gives us a good accuracy; however, we got lower TPR which might cause false prediction on patients who really have cancer. To avoid life-threatening mistake, I will choose 0.03 to be the best thresholds.